

Operating Systems Practice

Scheduler


Eunji Lee

(ejlee@ssu.ac.kr)

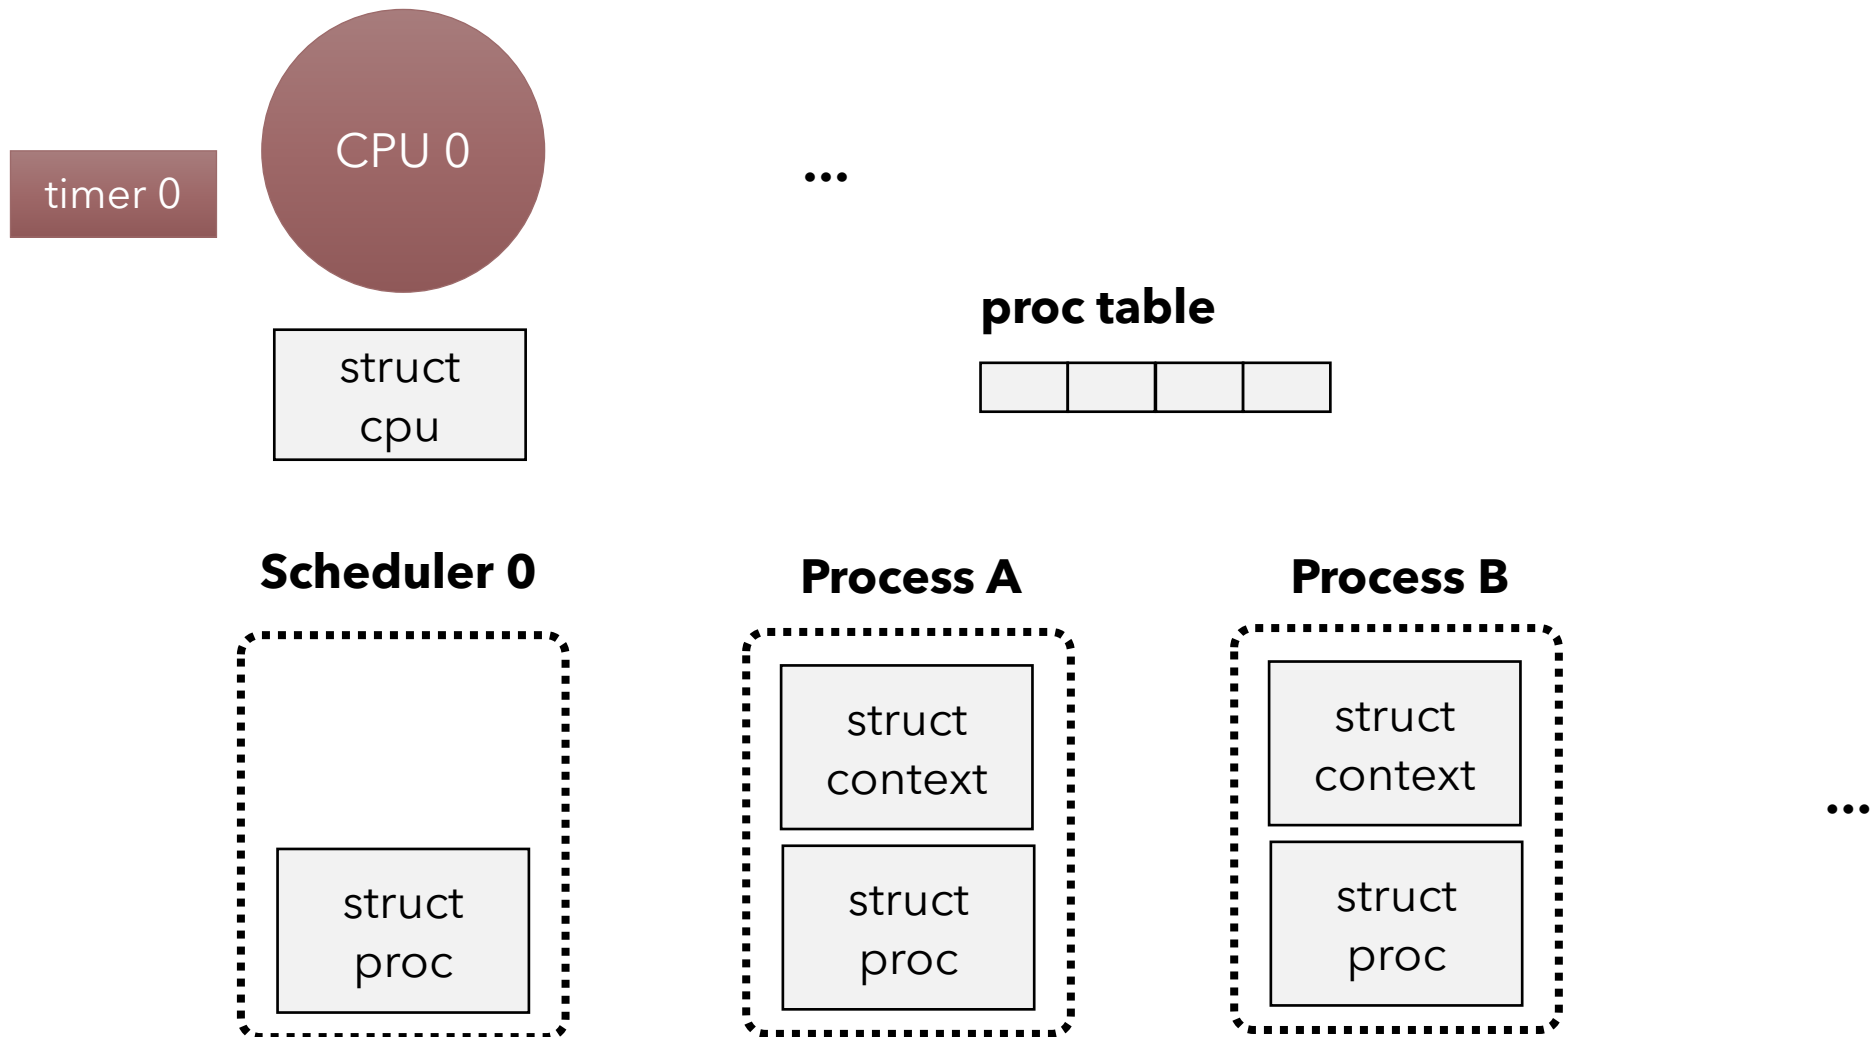


XV6 Scheduler – Round Robin

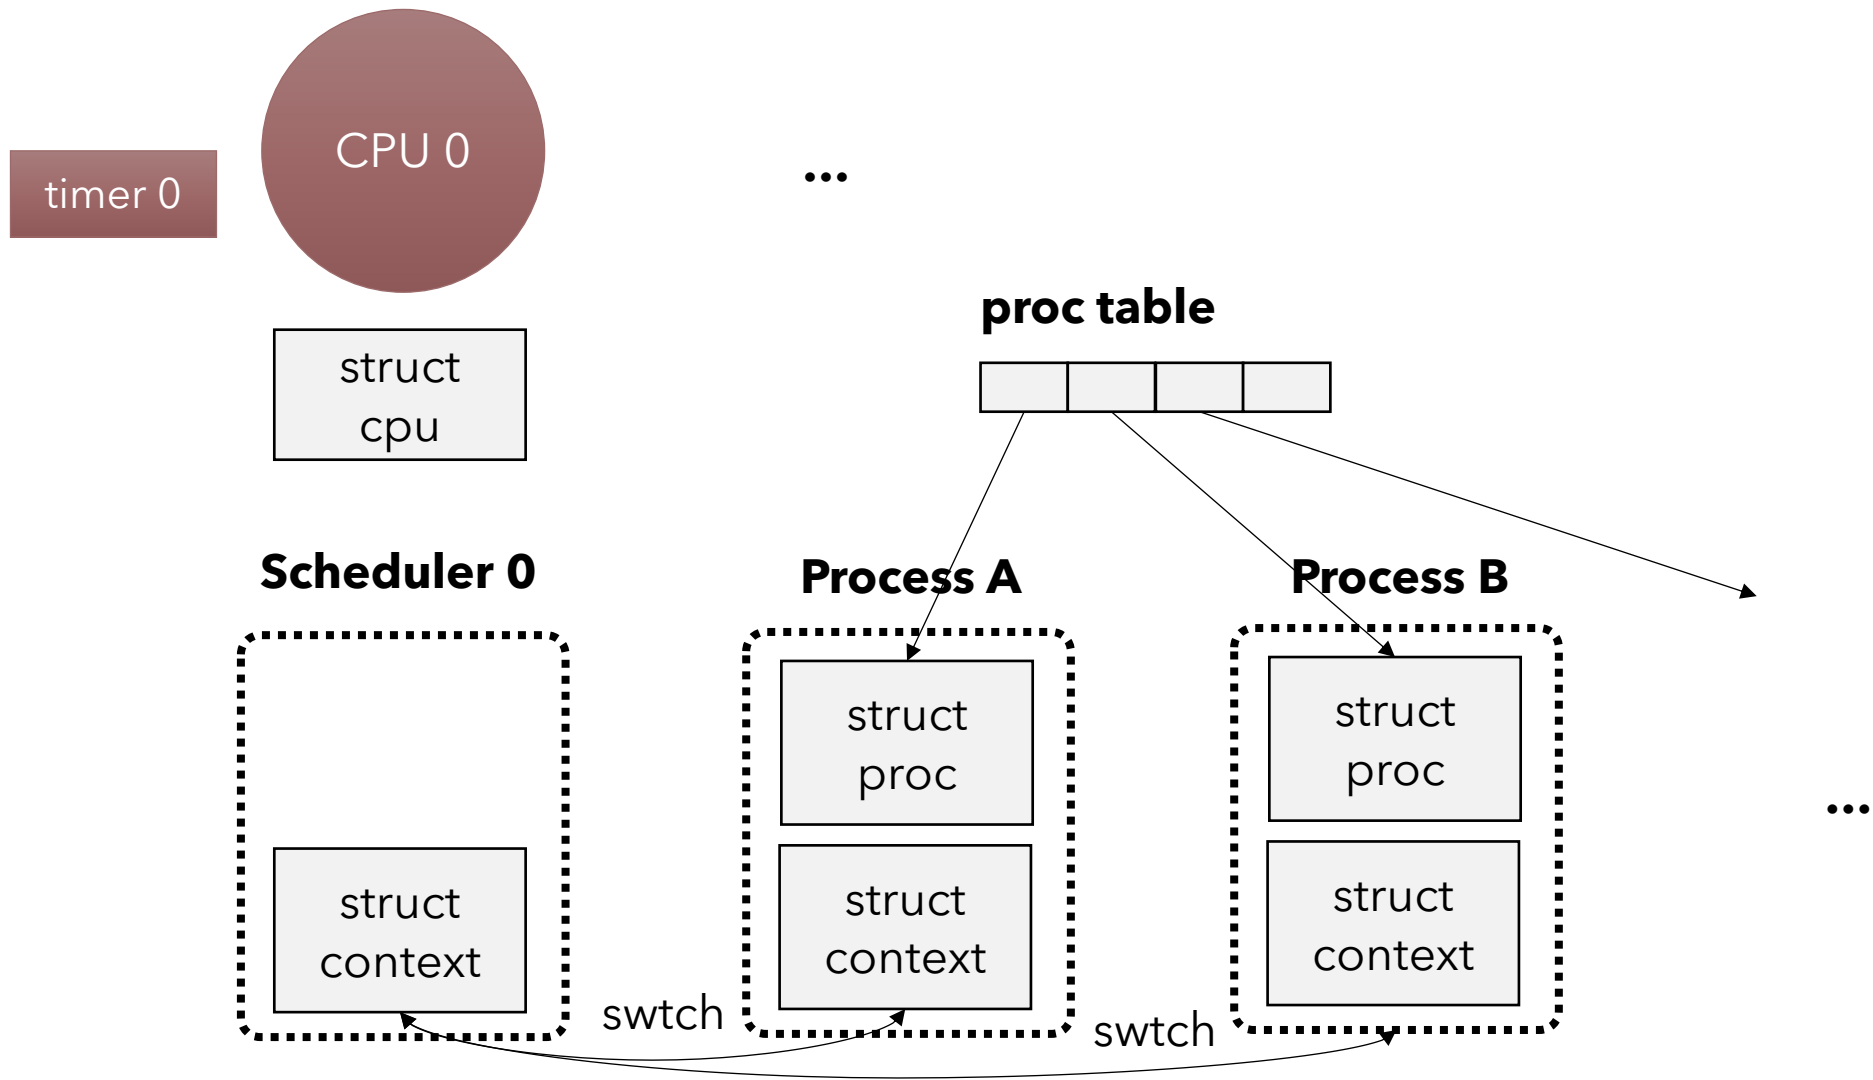
- Timer's interrupt request (IRQ) enforces a yield of a CPU
- A "RUNNABLE" process is chosen to be run in a round-robin manner



XV6 Data Structures for Scheduling



XV6 Data Structures for Scheduling



Process

- proc.h
- procstate
- struct proc

```
35 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
36
37 // Per-process state
38 struct proc {
39     uint sz;                // Size of process memory (bytes)
40     pde_t* pgdir;           // Page table
41     char *kstack;           // Bottom of kernel stack for this process
42     enum procstate state;    // Process state
43     int pid;                // Process ID
44     struct proc *parent;     // Parent process
45     struct trapframe *tf;    // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan;              // If non-zero, sleeping on chan
48     int killed;              // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd;       // Current directory
51     char name[16];           // Process name (debugging)
52     int priority;            // Process priority
53 };
```

Proc State

- UNUSED: Not used
- EMBRYO: Newly allocated (not ready for running yet)
- SLEEPING: Waiting for I/O, child process, or time
- RUNNABLE: Ready to run
- RUNNING: Running on CPU
- ZOMBIE: Exited

Process

- proc.h
- struct context

```
16 //PAGEBREAK: 17
17 // Saved registers for kernel context switches.
18 // Don't need to save all the segment registers (%cs, etc),
19 // because they are constant across kernel contexts.
20 // Don't need to save %eax, %ecx, %edx, because the
21 // x86 convention is that the caller has saved them.
22 // Contexts are stored at the bottom of the stack they
23 // describe; the stack pointer is the address of the context.
24 // The layout of the context matches the layout of the stack in switch.S
25 // at the "Switch stacks" comment. Switch doesn't save eip explicitly,
26 // but it is on the stack and allocproc() manipulates it.
27 struct context {
28     uint edi;
29     uint esi;
30     uint ebx;
31     uint ebp;
32     uint eip;
33 };
```

Scheduler

- proc.h
- struct cpu

```
1 // Per-CPU state
2 struct cpu {
3     uchar apicid;                // Local APIC ID
4     struct context *scheduler;  // swtch() here to enter scheduler
5     struct taskstate ts;        // Used by x86 to find stack for interrupt
6     struct segdesc gdt[NSEGS];  // x86 global descriptor table
7     volatile uint started;      // Has the CPU started?
8     int ncli;                   // Depth of pushcli nesting.
9     int intena;                 // Were interrupts enabled before pushcli?
10    struct proc *proc;          // The process running on this cpu or null
11 };
12
13 extern struct cpu cpus[NCPU];
14 extern int ncpu;
```

Scheduler

main.c

```
17 int
18 main(void)
19 {
20     kinit1(end, P2V(4*1024*1024)); // phys page allocator
21     kvmalloc(); // kernel page table
22     mpinit(); // detect other processors
23     lapicinit(); // interrupt controller
24     seginit(); // segment descriptors
25     picinit(); // disable pic
26     ioapicinit(); // another interrupt controller
27     consoleinit(); // console hardware
28     uartinit(); // serial port
29     pinit(); // process table
30     tvinit(); // trap vectors
31     binit(); // buffer cache
32     fileinit(); // file table
33     ideinit(); // disk
34     startothers(); // start other processors
35     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after kinit1
36     userinit(); // first user process
37     mpmain(); // finish this processor's setup
38 }
```

```
50 // Common CPU setup code.
51 static void
52 mpmain(void)
53 {
54     cprintf("cpu%d: starting %d\n", cpuid(), cpuid());
55     idtinit(); // load idt register
56     xchg(&(mycpu()->started), 1); // tell startothers() we're up
57     scheduler(); // start running processes
58 }
```

proc.c

```
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor.
331         sti();
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock);
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchuvm(p);
344             p->state = RUNNING;
345
346             switch(&(c->scheduler), p->context);
347             switchkvm();
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354     }
355 }
356 }
```

Start to execute
chosen process

When to Schedule

- `exit()`, `sleep()`
- timer interrupt (`yield()`)


trap.c

```
36 void
37 trap(struct trapframe *tf)
38 {
```

```
103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105 if(myproc() && myproc()->state == RUNNING &&
106     tf->trapno == T_IRQ0+IRQ_TIMER)
107     yield();
108
```

How Scheduler works

```
365 void
366 sched(void)
367 {
368     int intena;
369     struct proc *p = myproc();
370
371     if(!holding(&ptable.lock))
372         panic("sched ptable.lock");
373     if(mycpu()->ncli != 1)
374         panic("sched locks");
375     if(p->state == RUNNING)
376         panic("sched running");
377     if(readeflags() & FL_IF)
378         panic("sched interruptible");
379     intena = mycpu()->intena;
380     swtch(&p->context, mycpu()->scheduler);
381     mycpu()->intena = intena;
382 }
383
384 // Give up the CPU for one scheduling round.
385 void
386 yield(void)
387 {
388     acquire(&ptable.lock); //DOC: yieldlock
389     myproc()->state = RUNNABLE;
390     sched();
391     release(&ptable.lock);
392 }
```



How Scheduler works

```
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor.
331         sti();
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock);
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchvm(p);
344             p->state = RUNNING;
345
346             swtch(&(c->scheduler), p->context);
347             switchkvm();
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354     }
355 }
356 }
```

Return from here



Project 2. Priority Scheduling

- Implement priority scheduler in xv6
 - The lower nice value, the higher priority
 - The highest priority process should be chosen for next running
 - Tiebreak:Arbitrary
- Scheduler runs only when a change occurs in process priorities
 - DO NOT call the scheduler on the timer interrupt
 - When a process calls fork(), the **nice value** of child process is set to 5.

test_sched.c

- Add test_sched.c
- ./test_sched

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5
6 int main(int argc, char** argv)
7 {
8     int pid;
9     int mypid;
10
11     // Change the priority of init processes.
12     setnice(1, 10);
13
14     // Change the priority of current processes.
15     setnice(getpid(), 2);
16
17     // Create a child process
18     pid = fork();
19
20     if(pid == 0) {
21         printf(1, "#### State 2 ####\n");
22     } else {
23         printf(1, "#### State 1 ####\n");
24
25         // Change the priority of parent process.
26         setnice(pid, 10);
27         wait(); // Yield CPU
28
29         printf(1, "#### State 3 ####\n");
30     }
31
32     mypid = getpid();
33     printf(1, "PID %d is finished\n", mypid);
34
35     exit();
36 }
37
```

```
$ test_sched
#### State 1 ####
#### State 2 ####
PID 4 is finished
#### State 3 ####
PID 3 is finished
```


Hand-in Procedures (1/2)

- Download template
 - <https://github.com/eunjicious/xv6-ssu.git> (pull or clone)
 - `tar xvzf xv6_ssusyscall.tar.gz`
- Rename directory
 - `mv xv6_ssusyscall xv6_ssusched`
- Add test_sched.c to your codes and modify Makefile properly
- Build with CPUS=1 flag
 - Makefile

```
ifndef CPUS
CPUS := 1
endif
```

Hand-in Procedures (2/2)

- Compress your code (ID: 20201234)
 - `$tar cvzf xv6_ssu_sched_20201234.tar.gz xv6_ssu_sched`
 - Please command `$make clean` before compressing
- Submit your `tar.gz` file through myclass.ssu.ac.kr
- NO DELAY is allowed !!
- PLEASE DO NOT COPY !!