

Operating Systems Practice

Operating System Interface

Eunji Lee

(ejlee@ssu.ac.kr)



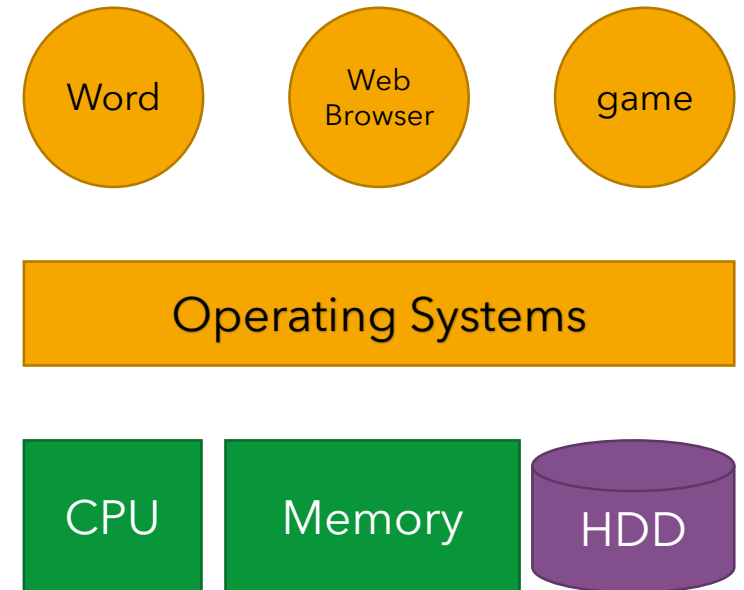


References

- Xv6-books
 - <https://pdos.csail.mit.edu/6.828/2018/xv6/book-rev10.pdf>
- Contents
 - **Ch.0: Operating system interface**
 - Ch.1: Operating system organization
 - Ch.2: Page tables
 - Ch.3: Traps, interrupts, and drivers
 - Ch.4: Locking
 - Ch.5: Scheduling
 - Ch.6: File system
 - Ch.7: Summary
 - Appendix A: PC hardware
 - Appendix B: The boot loader

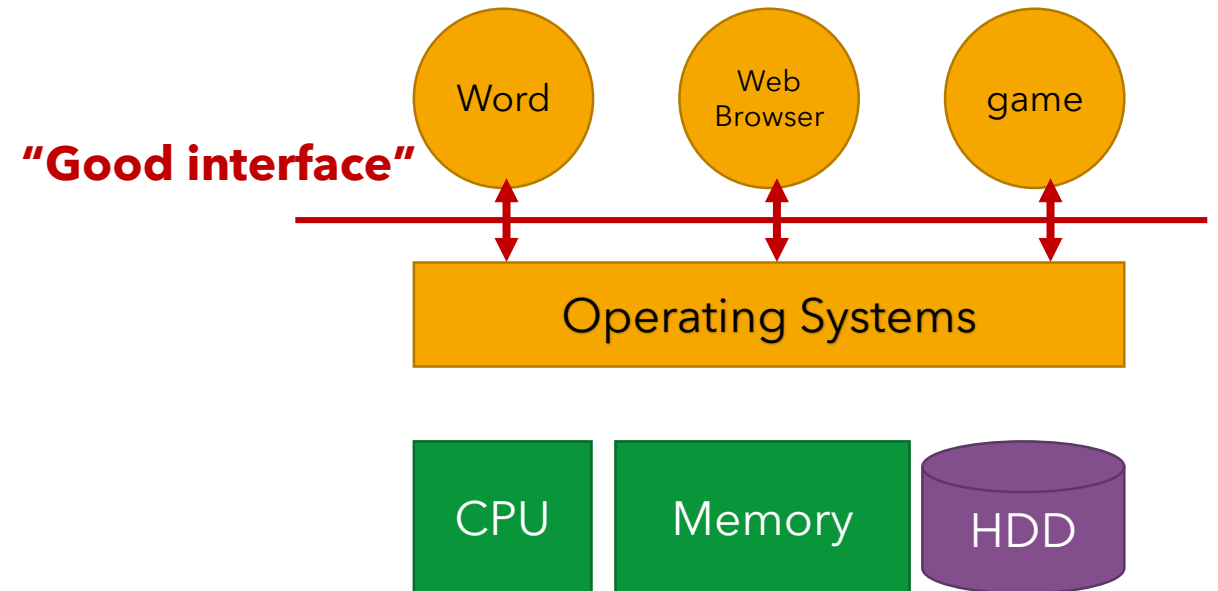
Job of Operating Systems

- **Share** a computer among multiple programs
- Manage and ***abstract*** the low-level hardware
- Run multiple programs at the same time
- Key ingredient: “Good interface”



Operating System Interfaces

- Requirements
 - Easy-to-use
 - Sophisticated features
- How to design interface?
 - Rely on **a few** mechanisms
 - **Combined** to provide **generality**



System call

- Operating system interface
- Invoked when a process needs to invoke a kernel service
- CPU's hardware protection
 - Ensure each process in user space can access only its own memory
 - When a user process invokes a system call, the hardware raises the **privilege level** and starts executing **a pre-arranged function (system call procedures)** in the kernel
- Services: Processes, memory, file system, pipes, etc.


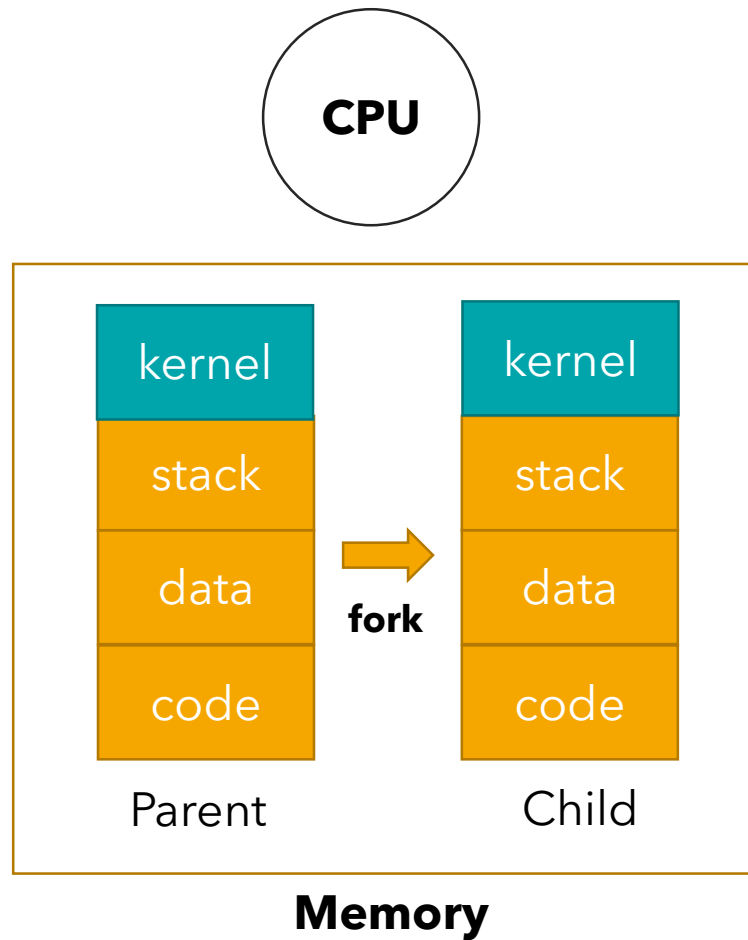


Figure 0-1. A kernel and two user processes.

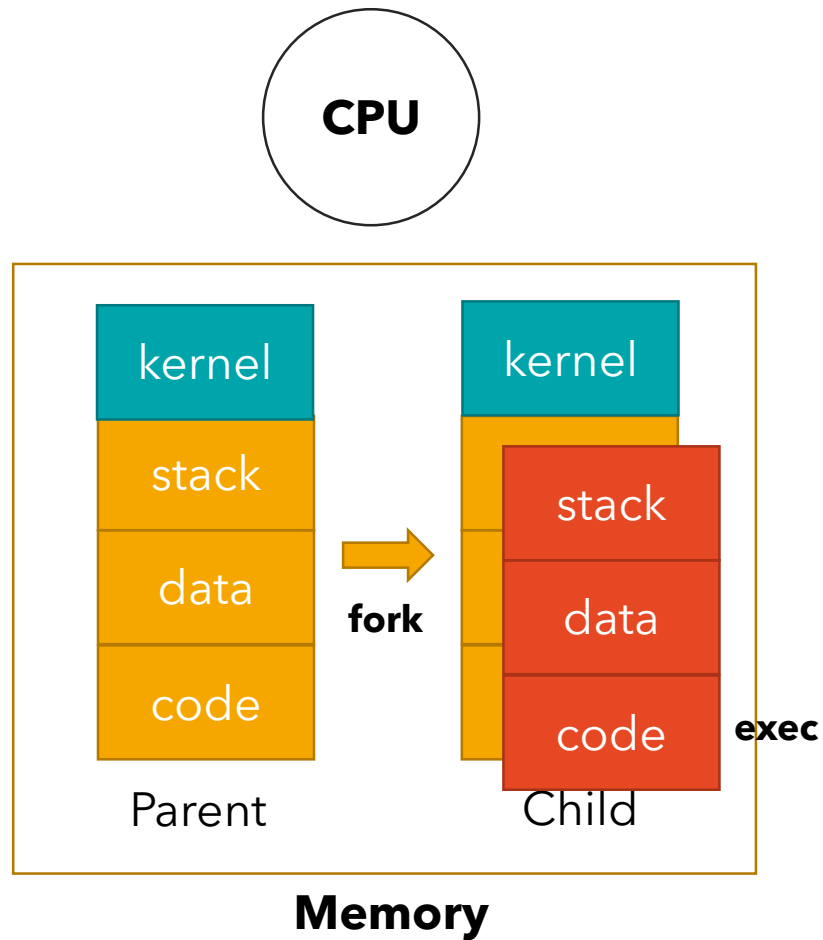
Processes and memory



```
// fork.c
int main()
{
    int pid = fork();

    if (pid > 0) {
        printf("parent: child=%d\n", pid);
        pid = wait();
        printf("child %d is done\n", pid);
    } else if (pid == 0) {
        printf("child: exiting\n");
    } else {
        printf("fork error\n");
    }
    return 0;
}
```

Processes and memory



```
// exec.c
int main()
{
    char *argv[3];
    argv[0] = "echo"
    argv[1] = "hello"
    argv[2] = 0;

    int pid = fork();
    ...
} else if (pid == 0) {
    printf("child: exiting\n");

    execve("/bin/echo", argv);
}
...
```

Processes and memory

- Shell

```
ejlee@ejlee-lecture:~/os20s$ ls -al
total 12
drwxrwxr-x  2 ejlee ejlee 4096 Feb 29 21:23 .
drwxr-xr-x 20 ejlee ejlee 4096 Feb 29 21:23 ..
-rw-rw-r--  1 ejlee ejlee   6 Feb 29 21:23 fork.c
ejlee@ejlee-lecture:~/os20s$
```


```
// sh.c

int main()
{
    while(1) {
        printf("$ ");
        getcmd(cmd);

        int pid = fork();
        if(pid > 0) {
            wait()
        } else if (pid == 0) {
            exec(cmd);
        }
    }
}
```


I/O and File descriptors

- File descriptor
 - A small integer representing a kernel object that a process may read from or write to
 - Abstract away the differences between files, pipes, and devices
 - Make them all look like streams of bytes
 - Kernel uses a file descriptor as an index into a per-process table



I/O and File descriptors

- cat

```
ejlee@ejlee-lecture:~/os20s$ ls
data.txt
ejlee@ejlee-lecture:~/os20s$ cat data.txt
This is a data file.
ejlee@ejlee-lecture:~/os20s$ █
```

I/O and File descriptors

- cat

```
// cat.c
void cat(int fd)
{
    int n;

    while((n = read(fd, buf, sizeof(buf))) > 0) {
        if (write(stdout, buf, n) != n) {
            printf("cat: write error\n");
            return;
        }
    }
    if(n < 0){
        printf("cat: read error\n");
        return;
    }
}
```

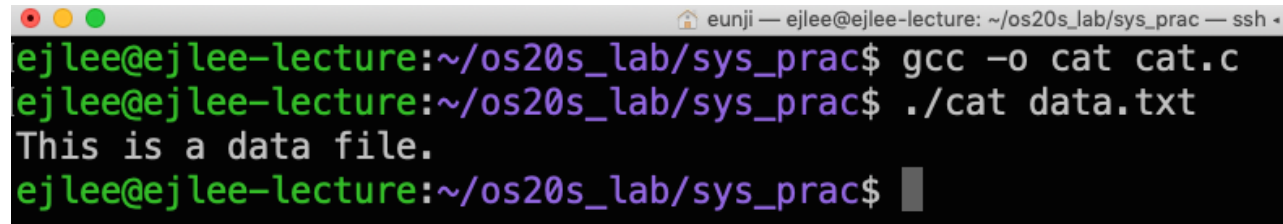
```
int main(int argc, char *argv[])
{
    int fd, i;

    if(argc <= 1){
        cat(0);
        exit();
    }

    for(i = 1; i < argc; i++){
        if((fd = open(argv[i], 0)) < 0){
            printf(1, "cat: cannot open %s\n", argv[i]);
            exit();
        }
        cat(fd);
        close(fd);
    }
    exit();
}
```

I/O and File descriptors

- cat

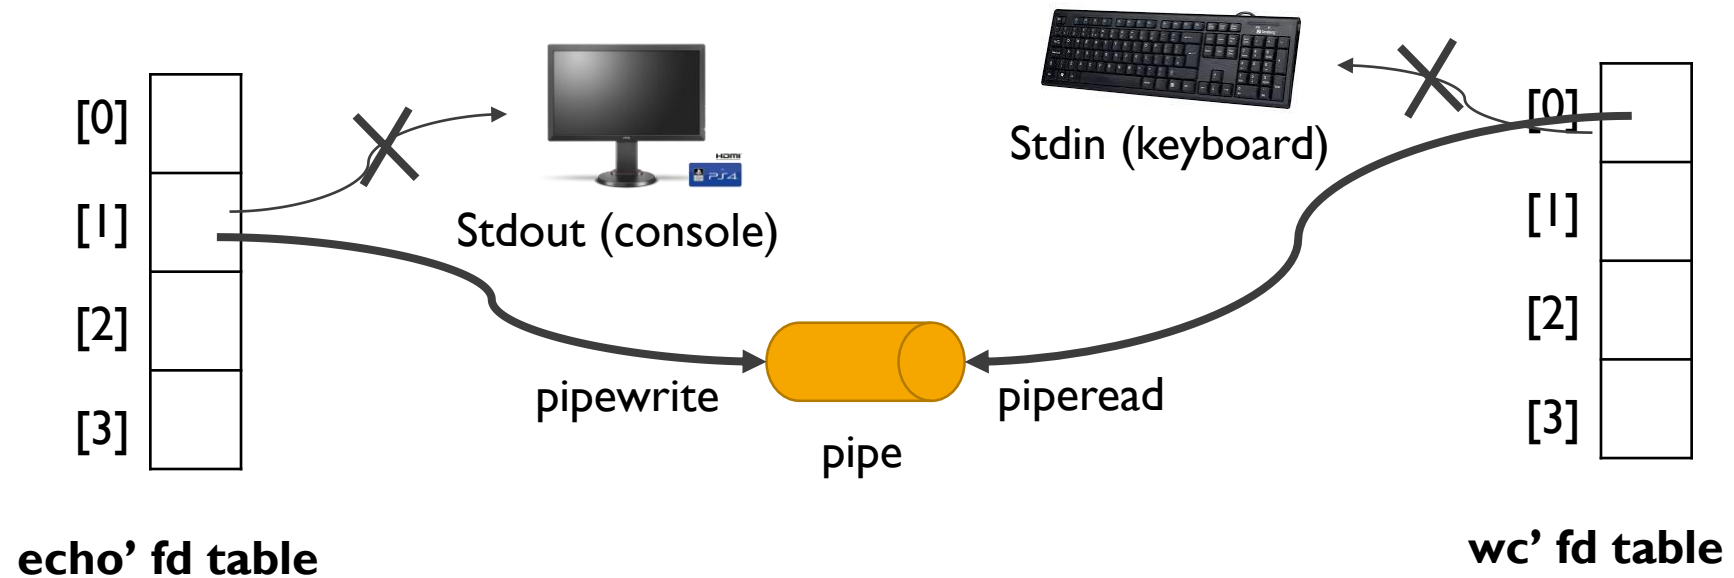


```
eunji — ejlee@ejlee-lecture: ~/os20s_lab/sys_prac — ssh
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$ gcc -o cat cat.c
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$ ./cat data.txt
This is a data file.
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$
```

Pipes

- A small kernel buffer exposed to processes as a pair of file descriptors
- Example

```
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$ echo "hello world" | wc
1      2      12
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$
```



Pipes


```
// pipe.c
int main()
{
    int p[2];
    char* argv[2];

    argv[0] = "wc";
    argv[1] = 0;

    pipe(p); // p[0]: readfd, p[1]: writefd
    if(fork() == 0) {
        close(0);
        dup(p[0]);
        close(p[0]);
        close(p[1]); // close unused write end
        execve("/usr/bin/wc", argv, NULL);
    } else {
        close(p[0]); // close unused read end
        write(p[1], "hello world\n", 12);
        close(p[1]);
    }
}
```

```
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$ ./pipe
1      2      12
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$
```

Before fork ..



Pipes

```
// pipe.c
int main()
{
    int p[2];
    char* argv[2];

    argv[0] = "wc";
    argv[1] = 0;

    pipe(p); // p[0]: readfd, p[1]: writefd
    if(fork() == 0) {
        close(0);
        dup(p[0]);
        close(p[0]);
        close(p[1]); // close unused write end
        execve("/usr/bin/wc", argv, NULL);
    } else {
        close(p[0]); // close unused read end
        write(p[1], "hello world\n", 12);
        close(p[1]);
    }
}
```

```
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$ ./pipe
1      2      12
ejlee@ejlee-lecture:~/os20s_lab/sys_prac$
```

