

Operating Systems Practice

Virtual Memory – Copy-on-Write

Eunji Lee

(ejlee@ssu.ac.kr)





Reference

- “xv6: a simple, Unix-like teaching operating system,” Chapter 3. Page tables

Paging Hardware

- RISC-V uses 39-bit virtual address
- A page table is stored in physical memory as a three-level tree
- Each pte has 44-bit PPN and 10-bit flag

mmu.h

```

134 // Page table/directory entry flags.
135 #define PTE_P 0x001 // Present
136 #define PTE_W 0x002 // Writeable
137 #define PTE_U 0x004 // User
138 #define PTE_PWT 0x008 // Write-Through
139 #define PTE_PCD 0x010 // Cache-Disable
140 #define PTE_A 0x020 // Accessed
141 #define PTE_D 0x040 // Dirty
142 #define PTE_PS 0x080 // Page Size
143 #define PTE_MBZ 0x180 // Bits must be zero
144

```

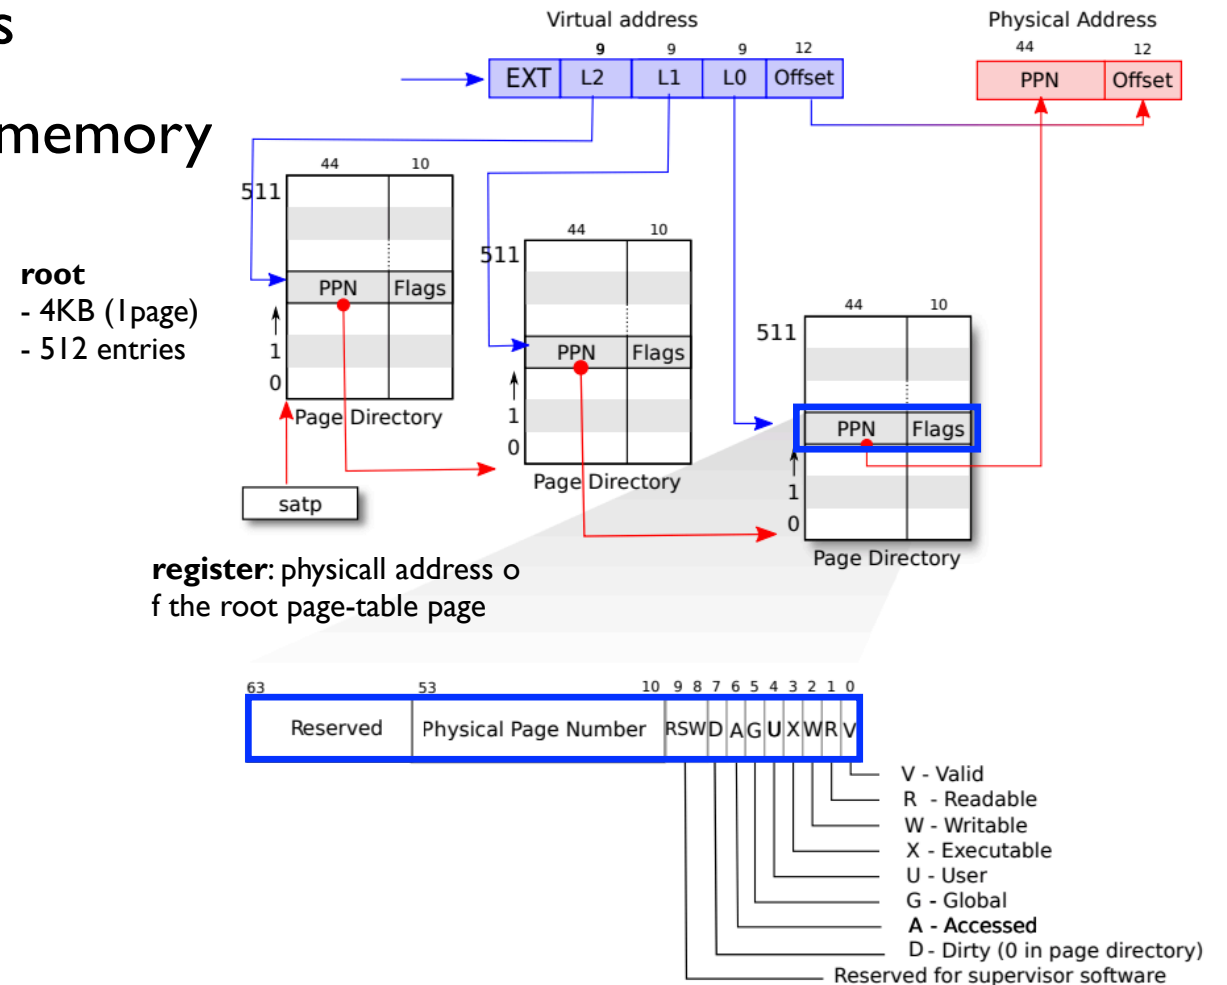


Figure 3.2: RISC-V page table hardware.

Kernel Address Space

- Direct mapped
- Only kernel stack pages are not direct-mapped
- guard page
 - Prevent problem caused by kernel stack overflow
 - guard page's PTE is invalid

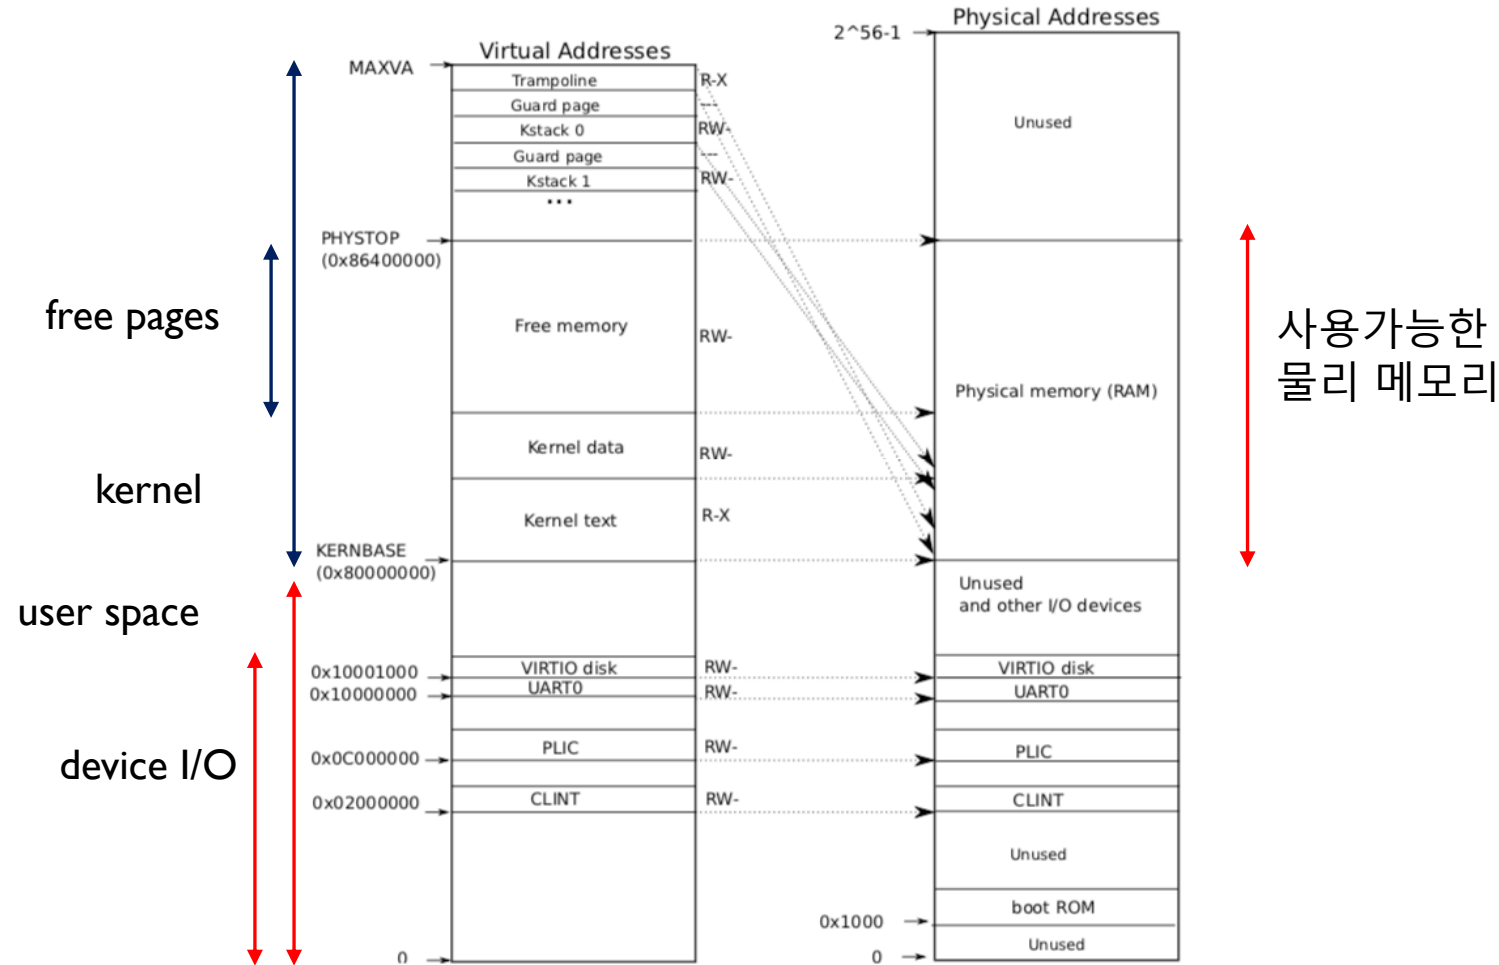


Figure 3.3: On the left, xv6's kernel address space. RWX refer to PTE read, write, and execute permissions. On the right, the RISC-V physical address space that xv6 expects to see.

Creating an address space

- vm.c : code for manipulate address spaces and page table

proc.c

```
int
fork(void)
{
    int i, pid;
    struct proc *np;
    struct proc *curproc = myproc();

    // Allocate process.
    if((np = allocproc()) == 0) { struct proc 할당
        return -1;
    }

    // Copy process state from proc. parent 의 address space 를 복사
    if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0) {
        kfree(np->kstack);
        np->kstack = 0;
        np->state = UNUSED;
        return -1;
    }
```

Creating an address space

vm.c

```
pde_t*
copyvm(pde_t *pgdir, uint sz)
{
    pde_t *d;
    pte_t *pte;
    uint pa, i, flags;
    char *mem;
```

pde_t *pgdir은 parent process의 page table
pde_t *d는 새로운 child process를 위한 page table

```
    if((d = setupkvm()) == 0)    kernel stack 할당
        return 0;
    for(i = 0; i < sz; i += PGSIZE){ // parent address space size = sz
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0) // walkpgdir: virtual address 에 대한 pte 찾기
            panic("copyvm: pte should exist");
        if(!(*pte & PTE_P))
            panic("copyvm: page not present");
        pa = PTE_ADDR(*pte);
        flags = PTE_FLAGS(*pte);
        if((mem = kalloc()) == 0)
            goto bad;
        memmove(mem, (char*) P2V(pa), PGSIZE);
        if(mappages(d, (void*) i, PGSIZE, V2P(mem), flags) < 0)
            goto bad;
    }
    return d;

bad:
    freevm(d);
    return 0;
}
```

page table entry에서 user virtual address에 대한 physical address 와 flag 를 읽어냄.

Creating an address space

vm.c

```
pde_t*
copyvm(pde_t *pgdir, uint sz)
{
    pde_t *d;
    pte_t *pte;
    uint pa, i, flags;
    char *mem;
```

```
    if((d = setupkvm()) == 0)
        return 0;
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
            panic("copyvm: pte should exist");
        if(!(*pte & PTE_P))
            panic("copyvm: page not present");
        pa = PTE_ADDR(*pte);
        flags = PTE_FLAGS(*pte);
```

```
        if((mem = kalloc()) == 0) physical page 하나를 얻어옴.
```

```
        goto bad;
```

```
        memmove(mem, (char*) P2V(pa), PGSIZE); parant 의 page 를 child 가 새로 할당받은 page 로 복사
```

```
        if(mappages(d, (void*) i, PGSIZE, V2P(mem), flags) < 0) 새로운 child 의 page 를 page table 에 저장
```

```
        goto bad;
```

```
    }
```

```
    return d;
```

d: child 의 page table

i: mapping 이 필요한 virtual address

V2P(mem) : physical address

Flags

```
bad:
```

```
    freevm(d);
```

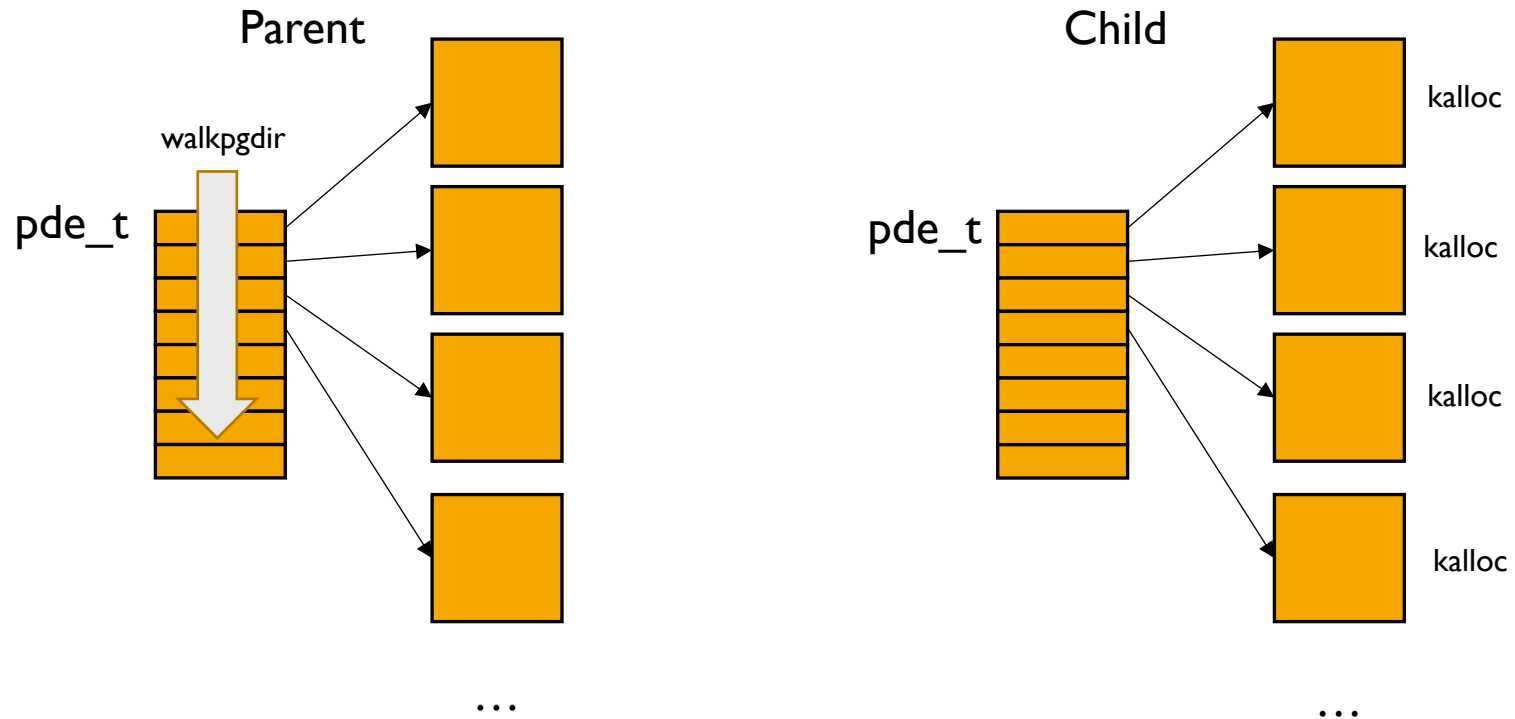
```
    return 0;
```

```
}
```

pde_t *pgdir은 parent process의 page table
pde_t *d는 새로운 child process를 위한 page table

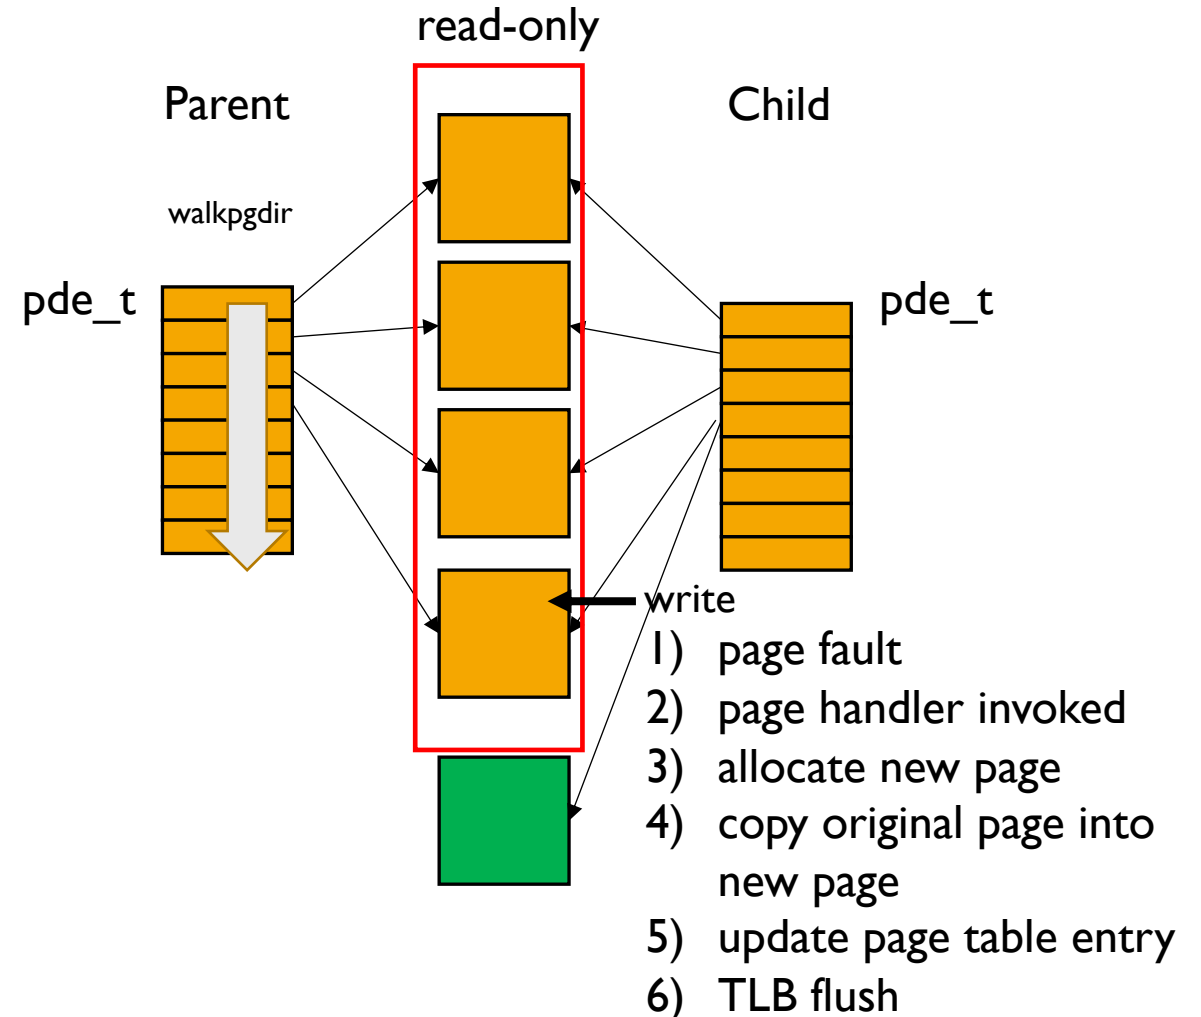
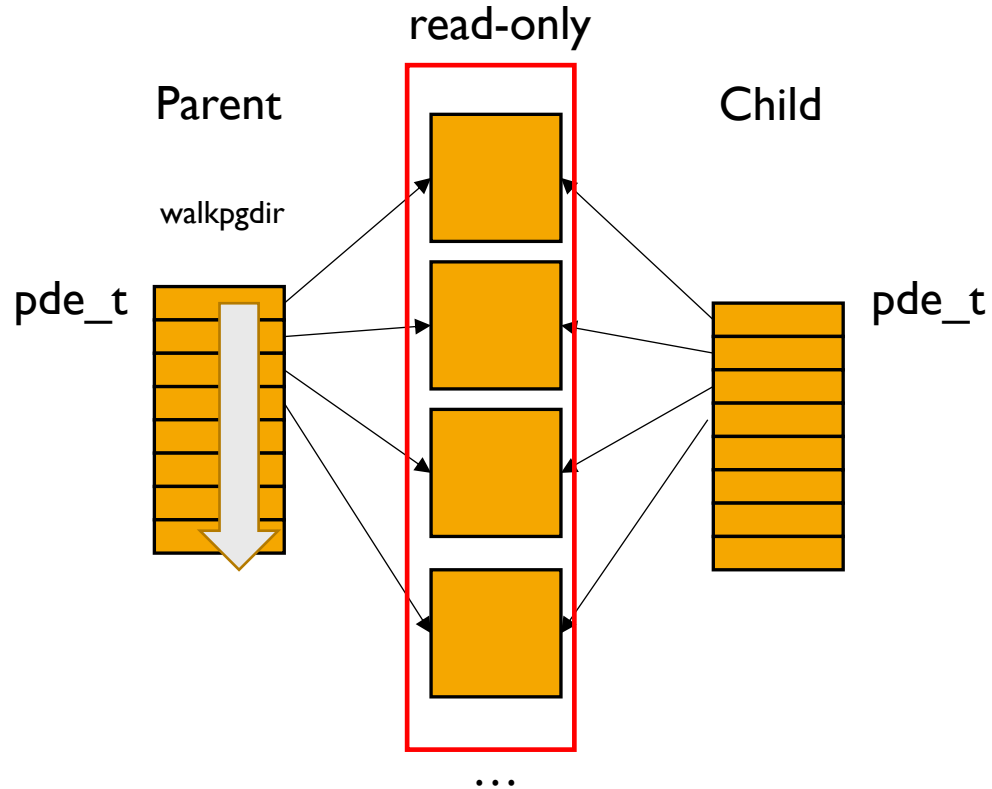
Creating an address space

- Duplication (copyuvm() in vm.c)



Creating an address space

- Copy-on-write (copyuvm() in vm.c)



Modification – Reference Counter

- Add reference counter and associated functions
 - kalloc.c

```
38 void
39 kinit1(void *vstart, void *vend)
40 {
41 #ifdef COW
42   initlock(&pmem.lock, "pmemlock");
43 #endif
44   initlock(&kmem.lock, "kmem");
45   kmem.use_lock = 0;
46   freerange(vstart, vend);
47 }
48
49 void
50 kinit2(void *vstart, void *vend)
51 {
52 #ifdef COW
53   memset(&pmem.ref, 1, sizeof(uint) * (PHYSTOP >> PGSHIFT));
54   acquire(&pmem.lock);
55   pmem.numfreepages = 0;
56   release(&pmem.lock);
57 #endif
58   freerange(vstart, vend);
59   kmem.use_lock = 1;
60 }
61
```

```
15 #ifdef COW
16 struct {
17   struct spinlock lock;
18   int numfreepages;
19   uint ref[PHYSTOP >> PGSHIFT];
20 } pmem;
21 #endif
```

```
98 //PAGEBREAK: 21
99 // Free the page of physical memory pointed at by v,
100 // which normally should have been returned by a
101 // call to kalloc(). (The exception is when
102 // initializing the allocator; see kinit above.)
103 void
104 kfree(char *v)
105 {
106   struct run *r;
107
108   if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
109     panic("kfree");
110
111   // Fill with junk to catch dangling refs.
112   memset(v, 1, PGSIZE);
113
114   if(kmem.use_lock)
115     acquire(&kmem.lock);
116   r = (struct run*)v;
117   r->next = kmem.freelist;
118   kmem.freelist = r;
119   if(kmem.use_lock)
120     release(&kmem.lock);
121 #ifdef COW
122   if(pmem.use_lock)
123     acquire(&pmem.lock);
124
125   pmem.numfreepages++;
126
127   if(pmem.use_lock)
128     release(&pmem.lock);
129 #endif
130 }
```

```
135 char*
136 kalloc(void)
137 {
138   struct run *r;
139
140   if(kmem.use_lock)
141     acquire(&kmem.lock);
142   r = kmem.freelist;
143   if(r){
144     kmem.freelist = r->next;
145
146 #ifdef COW
147   if(pmem.use_lock)
148     acquire(&pmem.lock);
149
150   pmem.numfreepages--;
151
152   if(pmem.use_lock)
153     release(&pmem.lock);
154 #endif
155 }
156   if(kmem.use_lock)
157     release(&kmem.lock);
158
159   return (char*)r;
160 }
161
```

Modification – Reference Counter

- Add reference counter and associated functions
 - kalloc.c

```
70 #ifdef COW
71 // reference counter APIs
72 int freemem()
73 {
74     return 0;
75 }
76
77 uint
78 get_ref(uint pa)
79 {
80     return 0;
81 }
82
83 void
84 inc_ref(void)
85 {
86     return;
87 }
88
89 void
90 dec_ref(uint pa)
91 {
92     return;
93 }
94 #endif
```

Modification – Reference Counter

- Add reference counter and associated functions
 - mmu.h

```
87 #define PGSHIFT 12 // log2(PGSIZE)
88 #define PTXSHIFT 12 // offset of PTX in a linear address
89 #define PDXSHIFT 22 // offset of PDX in a linear address
```

- memlayout.h

```
1 // Memory layout
2
3 #define EXTMEM 0x100000 // Start of extended memory
4 #define PHYSTOP 0xE000000 // Top physical memory
5 #define DEVSPACE 0xFE000000 // Other devices are at high addresses
6
```

Modification – Reference Counter

- Add reference counter and associated functions in `kalloc.c`
 - `ref[PHYSTOP >> PGSHIFT]`: reference counter for physical memory pages
 - `uint get_ref(uint pa)`: read the reference count of `pa`
 - `void inc_ref(uint pa)`: increase the reference count of `pa`
 - `void dec_ref(uint pa)`: decrease the reference count of `pa`
- Increase reference counter
 - When allocating the physical page
 - When referencing the physical page
- Decrease reference counter
 - When de-allocating the physical page
 - When de-referencing the physical page
- Increase / decrease reference counter appropriately
 - `kalloc` / `kfree`

Modification – Sharing pages

- copyvm() in vm.c
 - DO NOT allocate a new page for child's address space
 - Install parent's address spaces to child's page table
 - Set the page NOT writable (flags) for both child and parent
 - Increase reference counter for shared page
 - TLB flush: lcr3(V2P(pgdir))

vm.c

```
313 // Given a parent process's page table, create a copy
314 // of it for a child.
315 pde_t*
316 copyvm(pde_t *pgdir, uint sz)
317 {
318     pde_t *d;
319     pte_t *pte;
320     uint pa, i, flags;
321     char *mem;
322
323     if((d = setupkvm()) == 0)
324         return 0;
325     for(i = 0; i < sz; i += PGSIZE){
326         if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
327             panic("copyvm: pte should exist");
328         if(!(*pte & PTE_P))
329             panic("copyvm: page not present");
330         pa = PTE_ADDR(*pte);
331         flags = PTE_FLAGS(*pte);
332         if((mem = kalloc()) == 0)
333             goto bad;
334         memmove(mem, (char*)P2V(pa), PGSIZE);
335         if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0) {
336             kfree(mem);
337             goto bad;
338         }
339     }
340     return d;
341
342 bad:
343     freevm(d);
344     return 0;
345 }
```


Modification – Page fault handler

- Add page fault handler in trap.c
- T_PGFLT occurs when memory access is invalid

trap.c

```
35 //PAGEBREAK: 41
36 void
37 trap(struct trapframe *tf)
38 {
39     if(tf->trapno == T_SYSCALL){
40         if(myproc()->killed)
41             exit();
42         myproc()->tf = tf;
43         syscall();
44         if(myproc()->killed)
45             exit();
46         return;
47     }
48
49     switch(tf->trapno){
50 #ifdef COW
51     case T_PGFLT:
52         page_fault();
53         break;
54 #endif
55     case T_IRQ0 + IRQ_TIMER:
56         if(cpuid() == 0){
57             acquire(&tickslock);
58             ticks++;
59             wakeup(&ticks);
60             release(&tickslock);
61         }
62     }
```

Modification – Page fault handler

- `void page_fault(void)`
 - `rcr2()`: return virtual address incurring page fault
 - check if the virtual address is valid
 - locate page table entry for the virtual address
 - `pte_t`
 - `walkpgdir()`
 - check the reference count of the physical address corresponding to the virtual address
 - `get_refcounter()`
 - Perform copy-on-write
 - if reference counter > 2: allocate new page & copy
 - `kalloc(), memmove()`
 - decrease reference count for original page
 - Make pages writable (update page table entry)
 - `lcr3(V2P(pgdir))`: TLB flush

vm.c

```
387
388 #ifdef COW
389 void
390 page_fault(void)
391 {
392     uint va = rcr2();
393
394     // fill this part
395
396
397     return 0;
398 }
399 #endif
400
401 //PAGEBREAK!
402 // Blank page.
403 //PAGEBREAK!
404 // Blank page.
405 //PAGEBREAK!
406 // Blank page.
407
```

COW Test

- Three tests

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 int
6 main(int argc, char **argv)
7 {
8     int before, after;
9     int pid;
10
11     printf(1, "TEST1: ");
12
13     before = freemem();
14
15     pid = fork();
16     if(pid == 0){
17         after = freemem();
18         if(before - after == 68)
19             printf(1, "OK\n");
20         else
21             printf(1, "WRONG\n");
22         exit();
23     }
24     else{
25         wait();
26     }
27
28     exit();
29 }
```

cowtest1.c

1,1

All

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 int
6 main(int argc, char **argv)
7 {
8     int before, after;
9     int pid;
10
11     printf(1, "TEST2: ");
12
13     before = freemem();
14
15     pid = fork();
16     if(pid == 0){
17         exit();
18     }
19     else{
20         wait();
21     }
22
23     after = freemem();
24     if(before == after)
25         printf(1, "OK\n");
26     else
27         printf(1, "WRONG\n");
28
29     exit();
30 }
```

cowtest2.c

1,1

All

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 int global = 3;
6
7 int
8 main(int argc, char **argv)
9 {
10     int before, after;
11     int pid;
12
13     printf(1, "TEST3: ");
14
15     pid = fork();
16     if(pid == 0){
17         before = freemem();
18         global = 4;
19         after = freemem();
20         if(before - after == 1)
21             printf(1, "OK\n");
22         else
23             printf(1, "WRONG\n");
24         exit();
25     }
26     else{
27         wait();
28     }
29
30     exit();
31 }
```

cowtest3.c

1,1

All


COW Test

- All should be “OK”

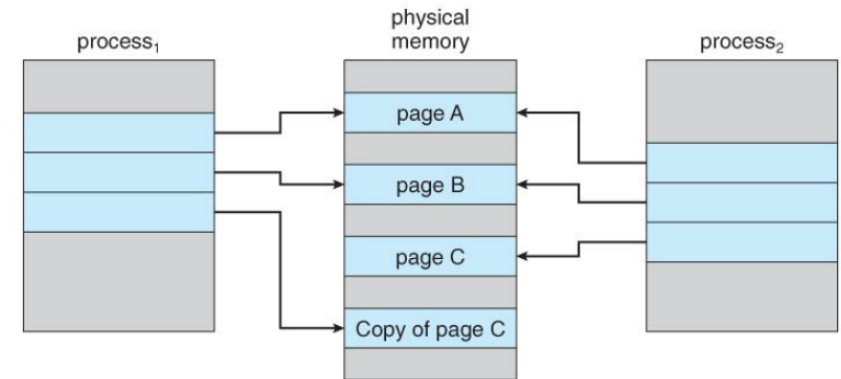
```
[ $ ls
.          1 1 512
..         1 1 512
README    2 2 2170
cat        2 3 13156
echo       2 4 12388
forktest   2 5 13068
grep        2 6 14600
init        2 7 12952
kill        2 8 12432
ln          2 9 12328
ls          2 10 14372
mkdir       2 11 12524
rm          2 12 12500
sh          2 13 21556
stressfs    2 14 13120
usertests   2 15 54636
wc          2 16 13604
zombie      2 17 12160
cowtest1    2 18 12592
cowtest2    2 19 12572
cowtest3    2 20 12716
console     3 21 0
[ $ cowtest1
TEST1: WRONG
[ $ cowtest2
TEST2: OK
[ $ cowtest3
TEST3: WRONG
$
```

Copy-on-Write

- When a process forks
 - Create shared mappings to the same page frames in physical page
 - Shared pages are protected as **read-only**
- When data is written to shared pages
 - Protection fault is generated
 - OS allocates new space in physical memory and directs the write to it
- **Reference counter for physical pages** is needed



Before process 1 modifies page C



After process 1 modifies page C